

Introduction into working with ocean simulation output



The Team



Ryan Patmore



Julia Rulent



Jonathan Coney

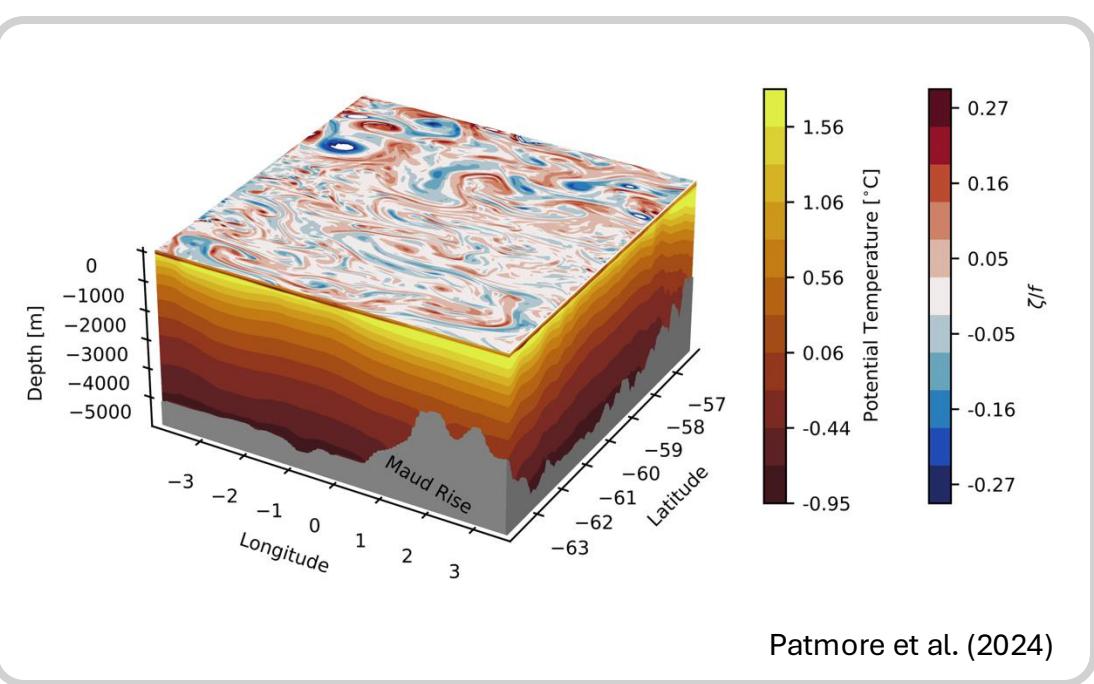


Bablu Sinha



The Aims

Hands-on appreciation for the nature of model data



Analysis methods for identifying tipping points

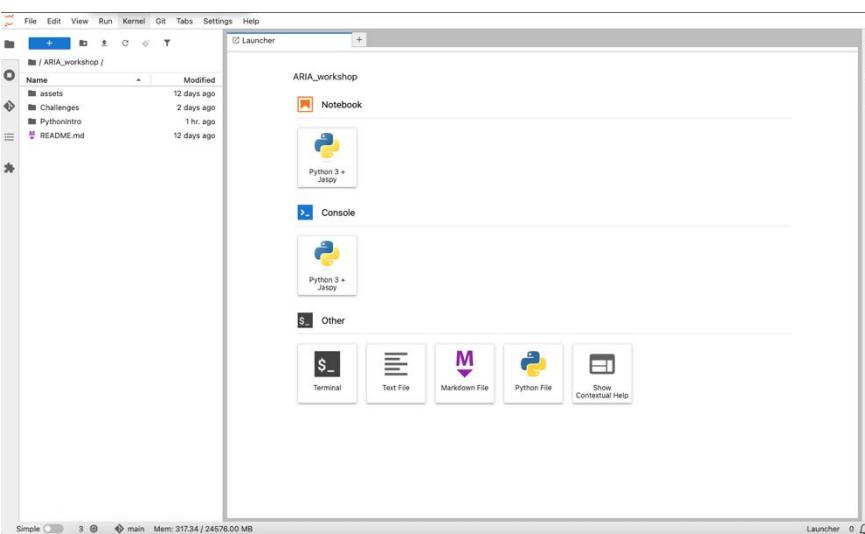


The Platforms

- JASMIN Jupyter notebooks for exploring the data
- Split into teams of 4/5 to tackle 2+ challenges
- Challenges outlined in ARIA_workshop repo

Notebook server

<https://notebooks.jasmin.ac.uk/>



ARIA_workshop Repository

https://github.com/NOC-MSM/ARIA_workshop

A screenshot of a GitHub repository page for 'ARIA_workshop'. The page shows a commit history with the following details:

Commit	Author	Message	Date
12dcb77	Julia Rulent	start challenge 2	4 days ago
main		1 Branch	0 Tags
13 Commits			

Below the commit history is a 'README' section with the title 'ARIA_workshop' and a sub-section 'Cloning this repository to the JASMIN Notebooks service' containing two steps:

1. Open up the JASMIN notebooks server. You might need to start your server if you haven't already.
2. Click the "git" icon on the left hand side of the window.

Instructions on landing page



The Content

- Reference content for using Python
- See Jonathan for advice ➔



ARIA_workshop Repository

Introductory Modules ➔

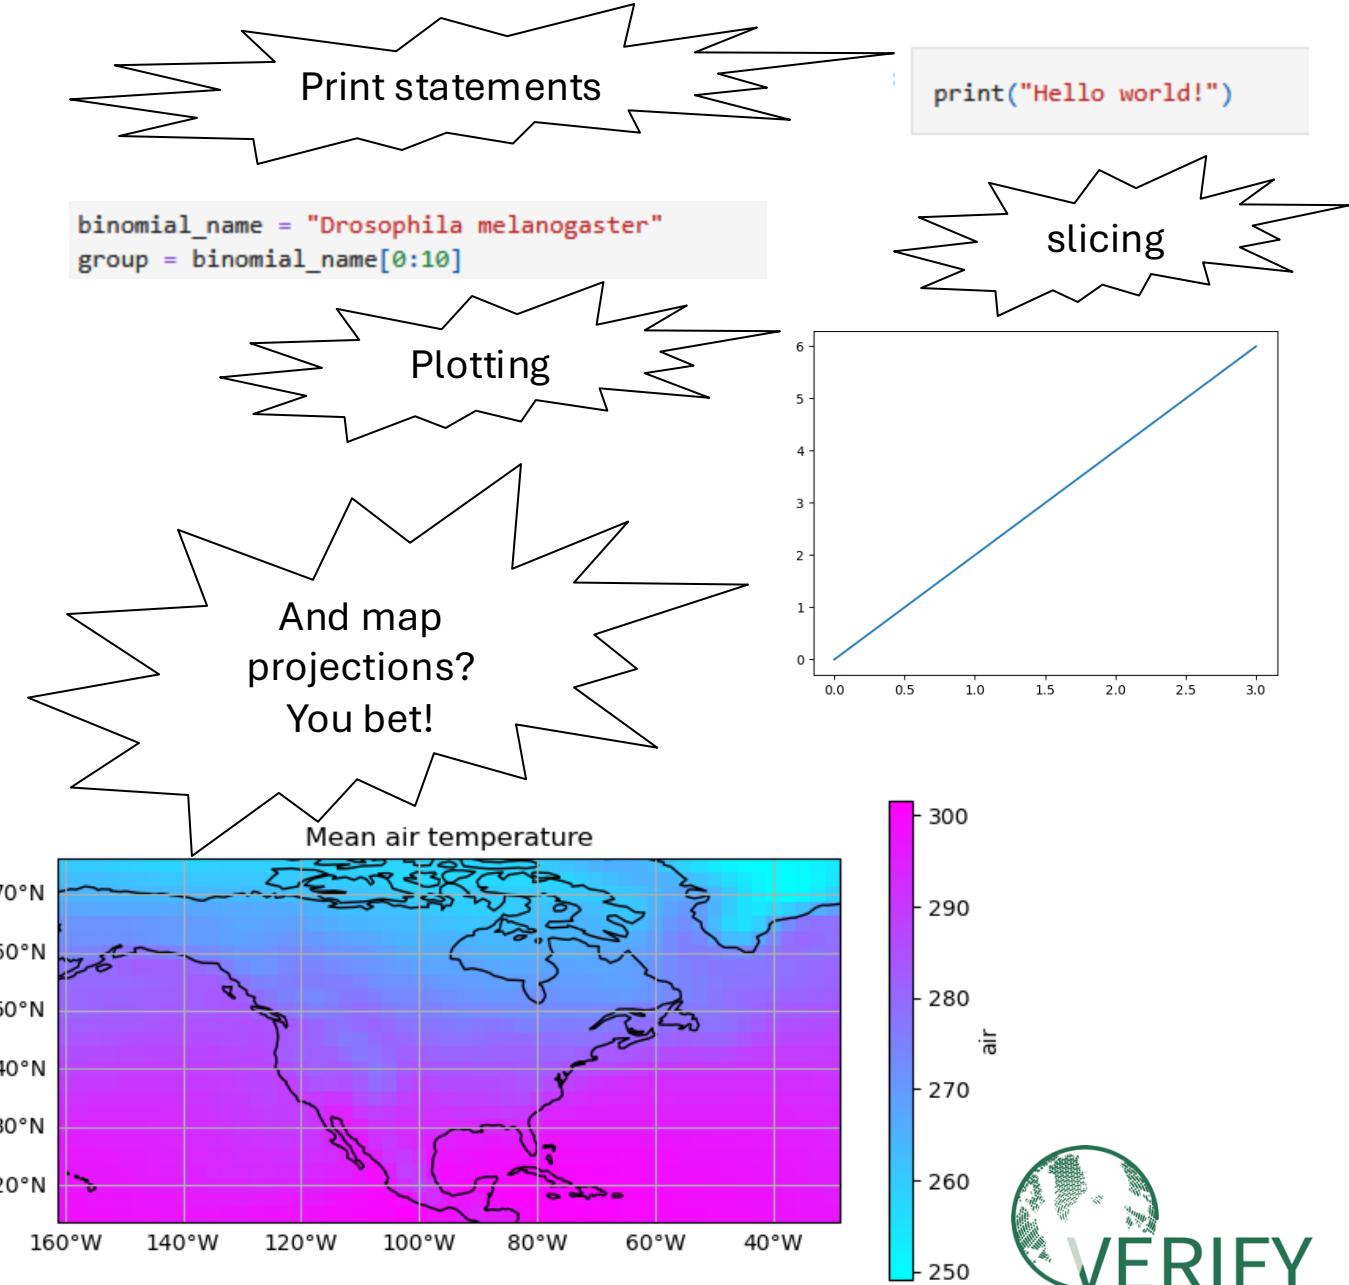
Julia Rulent	start challenge 2
Challenges	
PythonIntro	
assets	
README.md	
README	



PythonIntro introductory modules

- Thank you for filling in the poll with your Python ability.
- The Challenge notebooks assume knowledge of *some* Python, particularly familiarity with the libraries `xarray` and `matplotlib`.
- If you're fairly new to Python, unfamiliar with these libraries, or fancy refreshing your memory:
 - The two notebooks in the `PythonIntro/` directory are designed to be a whistle-stop tour of some basic Python (notebook 0) and an introduction to analysing and plotting data (notebook 1), in readiness for the Challenge notebooks.
 - Designed to be self-led, but if something doesn't make sense, **ask for help!**

It's all here:



The Content

Main content sits in Challenges directory

ARIA_workshop Repository

The group challenges



 Julia Rulent	start challenge 2
 Challenges	
 PythonIntro	
 assets	
 README.md	
 README	



The Content

Work through three challenges over 3x 1.5 hour sessions

ARIA_workshop Repository

Challenge 1 →

“Identify tipping points in the
CLASS model data”

Name
..
Challenge1.ipynb
Challenge2.ipynb
Challenge3_advanced.ipynb
README.md
README.md



The Content

Work through three challenges over 3x 1.5 hour sessions

ARIA_workshop Repository

Challenge 2 →

“Demonstrate how one might
devise an observational
campaign to capture this
tipping point”

Name
..
Challenge1.ipynb
Challenge2.ipynb
Challenge3_advanced.ipynb
README.md
README.md



The Content

Work through three challenges over 3x 1.5 hour sessions

ARIA_workshop Repository

“Derive tipping point metrics”

Challenge 3 (Advanced)



Name
..
Challenge1.ipynb
Challenge2.ipynb
Challenge3_advanced.ipynb
README.md
README.md



The Content

Are we all set up with the
JASMIN notebooks?



Challenge 1: Introduction

Grid Information

```
dom_cfg = xr.open_dataset(path + "domain_cfg.nc")
dom_cfg

Out[39]: <xarray.Dataset> Size: 8GB
Dimensions:      (y: 1207, x: 1442, nav_lev: 75, time_counter: 1, nlines: 67)
Coordinates:
  * nav_lev      (nav_lev) float32 300B 0.5058 1.556 ... 5.698e+03 5.902e+03
  * time_counter (time_counter) float64 8B 0.0
Dimensions without coordinates: y, x, nlines
Data variables: (12/43)
    nav_lon      (y, x) float32 7MB ...
    nav_lat      (y, x) float32 7MB ...
    jpiglo       int32 4B ...
    jjpglo       int32 4B ...
    jpkglo       int32 4B ...
    jperio       int32 4B ...
    ...
    ...
    bottom_level (time_counter, y, x) int32 7MB ...
    top_level    (time_counter, y, x) int32 7MB ...
    isf_draft    (time_counter, y, x) float64 14MB ...
    bathy_metry  (time_counter, y, x) float64 14MB ...
    namelist_cfg (nlines) |S102 7kB ...
    closea_mask   (y, x) float64 14MB ...
Attributes:
    DOMAIN_number_total:      1
    DOMAIN_number:            0
    DOMAIN_dimensions_ids:   [1 2]
    DOMAIN_size_global:      [1442 1207]
    DOMAIN_size_local:        [1442 1207]
    DOMAIN_position_first:   [1 1]
    DOMAIN_position_last:    [1442 1207]
    DOMAIN_halo_size_start:  [0 0]
    DOMAIN_halo_size_end:    [0 0]
    DOMAIN_type:              BOX
```



Challenge 1: Introduction

Depth and
time
coordinates

Grid Information

```
dom_cfg = xr.open_dataset(path + "domain_cfg.nc")
dom_cfg

Out[39]: <xarray.Dataset> Size: 8GB
Dimensions:      (y: 1207, x: 1442, nav_lev: 75, time_counter: 1, nlines: 67)
Coordinates:
  * nav_lev      (nav_lev) float32 300B 0.5058 1.556 ... 5.698e+03 5.902e+03
  * time_counter (time_counter) float64 8B 0.0
Dimensions without coordinates: y, x, nlines
Data variables: (12/43)
    nav_lon      (y, x) float32 7MB ...
    nav_lat      (y, x) float32 7MB ...
    jpiglo       int32 4B ...
    jjpglo       int32 4B ...
    jpkgl0       int32 4B ...
    jperio       int32 4B ...
    ...
    ...
    bottom_level (time_counter, y, x) int32 7MB ...
    top_level    (time_counter, y, x) int32 7MB ...
    isf_draft    (time_counter, y, x) float64 14MB ...
    bathy_metry  (time_counter, y, x) float64 14MB ...
    namelist_cfg (nlines) |S102 7kB ...
    closea_mask   (y, x) float64 14MB ...
Attributes:
    DOMAIN_number_total:      1
    DOMAIN_number:            0
    DOMAIN_dimensions_ids:   [1 2]
    DOMAIN_size_global:      [1442 1207]
    DOMAIN_size_local:       [1442 1207]
    DOMAIN_position_first:  [1 1]
    DOMAIN_position_last:   [1442 1207]
    DOMAIN_halo_size_start: [0 0]
    DOMAIN_halo_size_end:   [0 0]
    DOMAIN_type:             BOX
```



Challenge 1: Introduction

Latitude and Longitude

Grid Information

```
dom_cfg = xr.open_dataset(path + "domain_cfg.nc")
dom_cfg

Out[39]: <xarray.Dataset> Size: 8GB
Dimensions:      (y: 1207, x: 1442, nav_lev: 75, time_counter: 1, nlines: 67)
Coordinates:
  * nav_lev      (nav_lev) float32 300B 0.5058 1.556 ... 5.698e+03 5.902e+03
  * time_counter (time_counter) float64 8B 0.0
Dimensions without coordinates: y, x, nlines
Data variables: (12/43)
  nav_lon      (y, x) float32 7MB ...
  nav_lat      (y, x) float32 7MB ...
  jpiglo       int32 4B ...
  jjpglo       int32 4B ...
  jpkgl0       int32 4B ...
  jperio       int32 4B ...
  ...
  ...
  bottom_level (time_counter, y, x) int32 7MB ...
  top_level    (time_counter, y, x) int32 7MB ...
  isf_draft    (time_counter, y, x) float64 14MB ...
  bathy_metry  (time_counter, y, x) float64 14MB ...
  namelist_cfg (nlines) |S102 7kB ...
  closea_mask   (y, x) float64 14MB ...
Attributes:
  DOMAIN_number_total:      1
  DOMAIN_number:            0
  DOMAIN_dimensions_ids:   [1 2]
  DOMAIN_size_global:       [1442 1207]
  DOMAIN_size_local:        [1442 1207]
  DOMAIN_position_first:   [1 1]
  DOMAIN_position_last:    [1442 1207]
  DOMAIN_halo_size_start:  [0 0]
  DOMAIN_halo_size_end:    [0 0]
  DOMAIN_type:              BOX
```



Challenge 1: Introduction

Grid Information

The remainder is
general grid
information

```
dom_cfg = xr.open_dataset(path + "domain_cfg.nc")
dom_cfg

Out[39]: <xarray.Dataset> Size: 8GB
Dimensions:      (y: 1207, x: 1442, nav_lev: 75, time_counter: 1, nlines: 67)
Coordinates:
  * nav_lev      (nav_lev) float32 300B 0.5058 1.556 ... 5.698e+03 5.902e+03
  * time_counter (time_counter) float64 8B 0.0
Dimensions without coordinates: y, x, nlines
Data variables: (12/43)
    nav_lon      (y, x) float32 7MB ...
    nav_lat      (y, x) float32 7MB ...
    jpiglo       int32 4B ...
    jjpglo       int32 4B ...
    jpkgl0       int32 4B ...
    jperio       int32 4B ...
    ...
    ...
    bottom_level (time_counter, y, x) int32 7MB ...
    top_level    (time_counter, y, x) int32 7MB ...
    isf_draft    (time_counter, y, x) float64 14MB ...
    bathy_metry  (time_counter, y, x) float64 14MB ...
    namelist_cfg (nlines) |S102 7kB ...
    closea_mask   (y, x) float64 14MB ...
Attributes:
    DOMAIN_number_total:      1
    DOMAIN_number:            0
    DOMAIN_dimensions_ids:   [1 2]
    DOMAIN_size_global:       [1442 1207]
    DOMAIN_size_local:        [1442 1207]
    DOMAIN_position_first:   [1 1]
    DOMAIN_position_last:    [1442 1207]
    DOMAIN_halo_size_start:  [0 0]
    DOMAIN_halo_size_end:    [0 0]
    DOMAIN_type:              BOX
```



Challenge 1: Introduction

Grid Information

The remainder is
general grid
information

Let's delve
further...

```
dom_cfg = xr.open_dataset(path + "domain_cfg.nc")
dom_cfg

Out[39]: <xarray.Dataset> Size: 8GB
Dimensions:      (y: 1207, x: 1442, nav_lev: 75, time_counter: 1, nlines: 67)
Coordinates:
  * nav_lev      (nav_lev) float32 300B 0.5058 1.556 ... 5.698e+03 5.902e+03
  * time_counter (time_counter) float64 8B 0.0
Dimensions without coordinates: y, x, nlines
Data variables: (12/43)
    nav_lon      (y, x) float32 7MB ...
    nav_lat      (y, x) float32 7MB ...
    jpiglo       int32 4B ...
    jjpglo       int32 4B ...
    jpkgl0       int32 4B ...
    jperio       int32 4B ...
    ...
    ...
    bottom_level (time_counter, y, x) int32 7MB ...
    top_level    (time_counter, y, x) int32 7MB ...
    isf_draft    (time_counter, y, x) float64 14MB ...
    bathy_metry  (time_counter, y, x) float64 14MB ...
    namelist_cfg (nlines) |S102 7kB ...
    closea_mask   (y, x) float64 14MB ...
Attributes:
    DOMAIN_number_total:      1
    DOMAIN_number:            0
    DOMAIN_dimensions_ids:   [1 2]
    DOMAIN_size_global:       [1442 1207]
    DOMAIN_size_local:        [1442 1207]
    DOMAIN_position_first:   [1 1]
    DOMAIN_position_last:    [1442 1207]
    DOMAIN_halo_size_start:  [0 0]
    DOMAIN_halo_size_end:    [0 0]
    DOMAIN_type:              BOX
```



Challenge 1: Introduction

NEMO uses the commonly adopted Arakawa
C-Grid (Arakawa and Lamb, 1977)



Challenge 1: Introduction

NEMO uses the commonly adopted Arakawa C-Grid (Arakawa and Lamb, 1977)

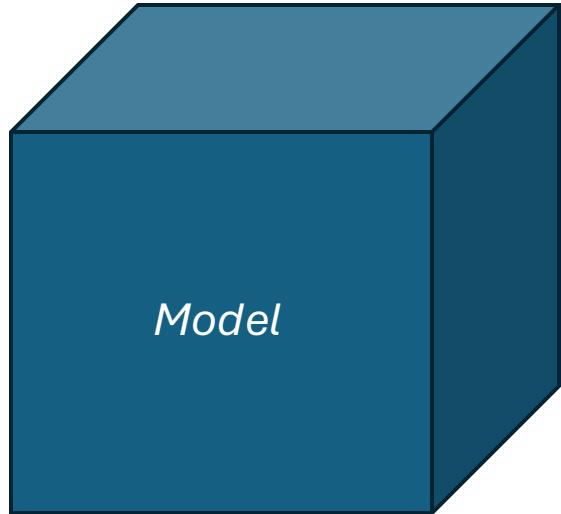
*It also runs
on Fortran!*



Challenge 1: Introduction

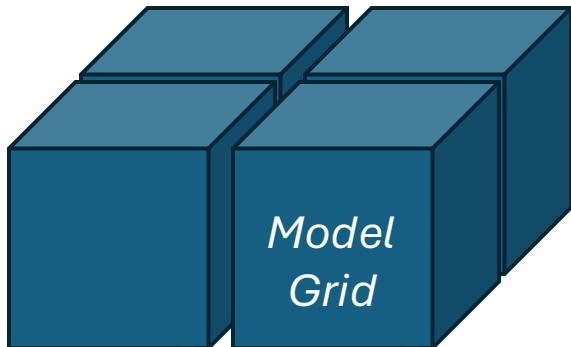
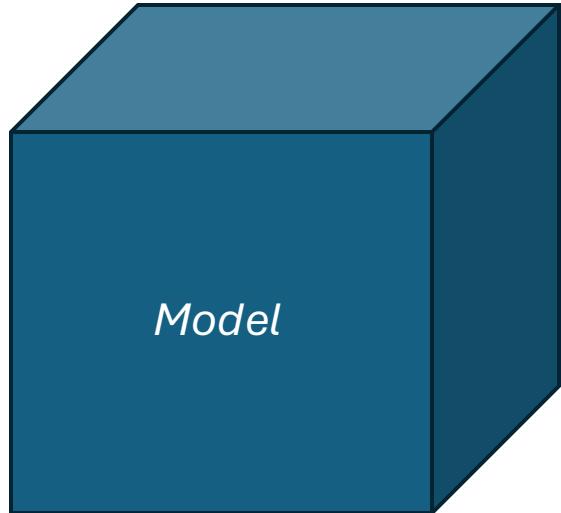
NEMO uses the commonly adopted Arakawa C-Grid (Arakawa and Lamb, 1977)

*It also runs
on Fortran!*



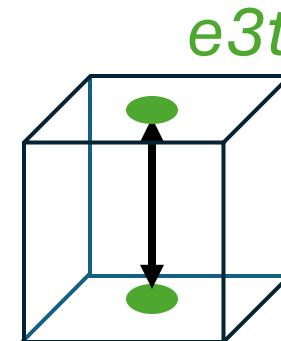
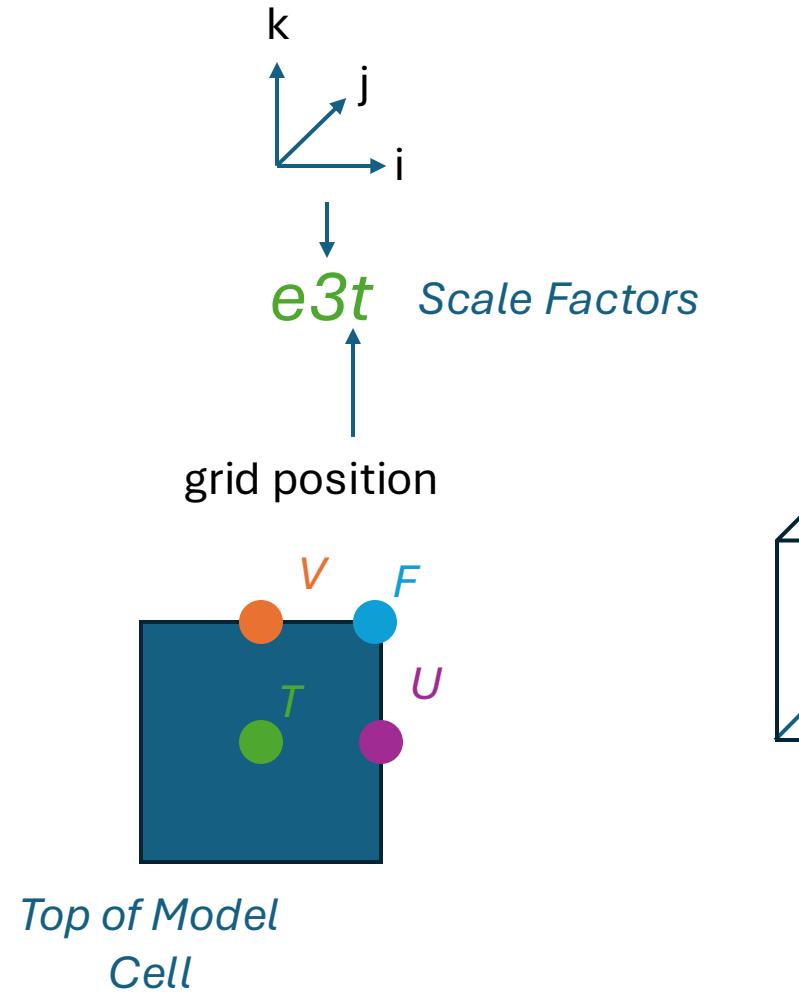
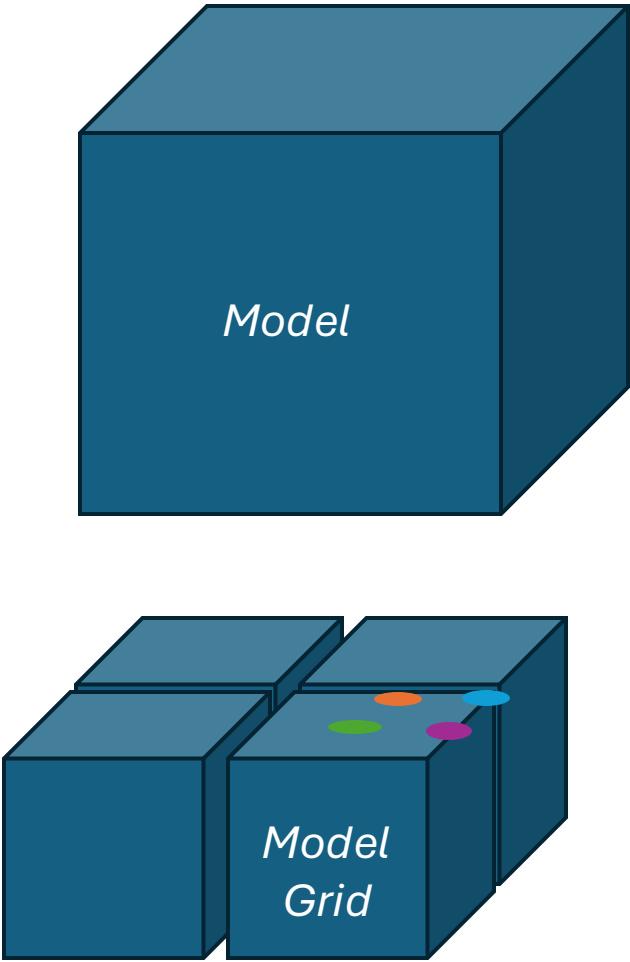
Challenge 1: Introduction

NEMO uses the commonly adopted Arakawa C-Grid (Arakawa and Lamb, 1977)



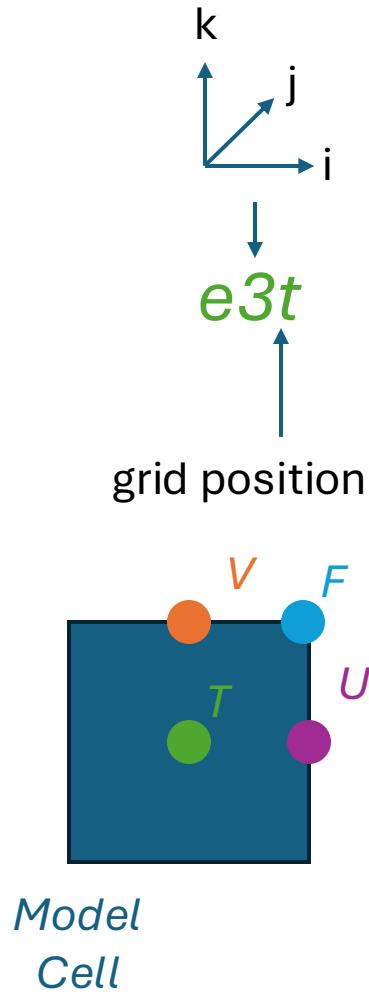
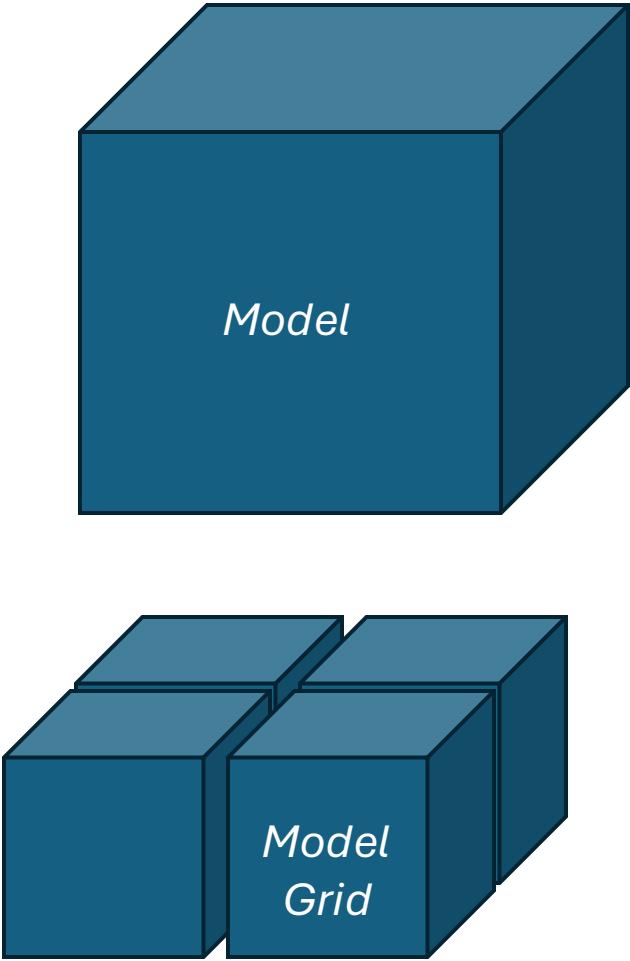
Challenge 1: Introduction

NEMO uses the commonly adopted Arakawa C-Grid (Arakawa and Lamb, 1977)



Challenge 1: Introduction

NEMO uses the commonly adopted Arakawa C-Grid (Arakawa and Lamb, 1977)



domain_cfg.nc

15]: Data variables:

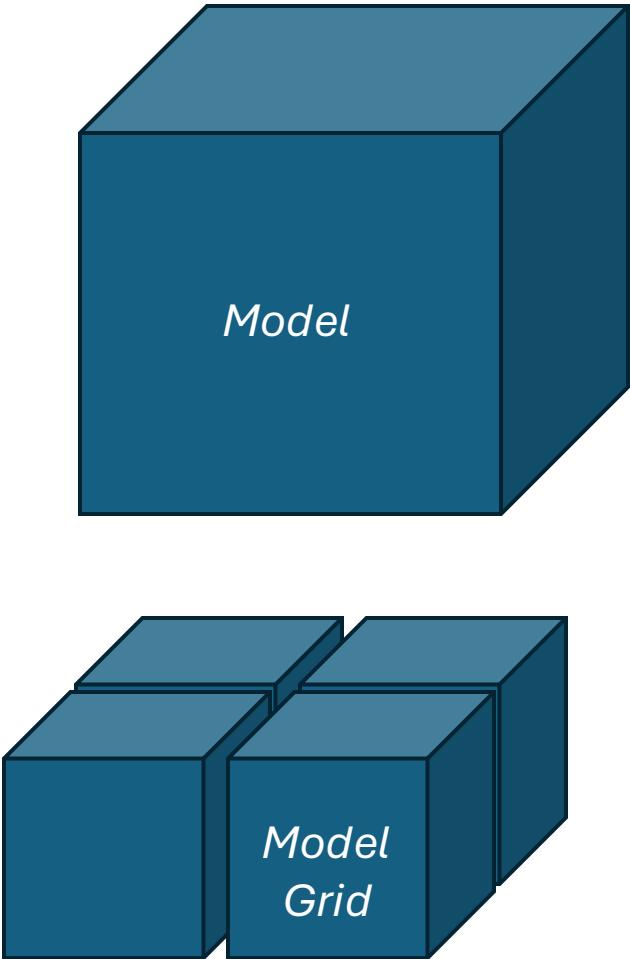
```
bottom_level    (t, y, x) int32 640kB ...
closea_mask     (y, x) float64 1MB ...
e1f              (t, y, x) float64 1MB ...
e1t              (t, y, x) float64 1MB ...
e1u              (t, y, x) float64 1MB ...
e1v              (t, y, x) float64 1MB ...
e2f              (t, y, x) float64 1MB ...
e2t              (t, y, x) float64 1MB ...
e2u              (t, y, x) float64 1MB ...
e2v              (t, y, x) float64 1MB ...
e3f_0            (t, z, y, x) float64 96MB ...
e3t_0            (t, z, y, x) float64 96MB ...
e3t_1d           (t, z) float64 600B ...
e3u_0            (t, z, y, x) float64 96MB ...
e3uw_0           (t, z, y, x) float64 96MB ...
e3v_0            (t, z, y, x) float64 96MB ...
e3vw_0           (t, z, y, x) float64 96MB ...
e3w_0            (t, z, y, x) float64 96MB ...
e3w_1d           (t, z) float64 600B ...
ff_f              (t, y, x) float64 1MB ...
ff_t              (t, y, x) float64 1MB ...
glamf             (t, y, x) float64 1MB ...
glamt             (t, y, x) float64 1MB ...
glamu             (t, y, x) float64 1MB ...
glamv             (t, y, x) float64 1MB ...
gphif             (t, y, x) float64 1MB ...
gphit             (t, y, x) float64 1MB ...
gphiu             (t, y, x) float64 1MB ...
gphiv             (t, y, x) float64 1MB ...
jperio            int32 4B ...
jpiglo            int32 4B ...
```

Scale Factors

Latitude and Longitude



Challenge 1: Introduction

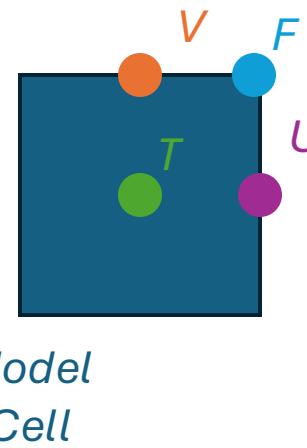


T-Grid Data

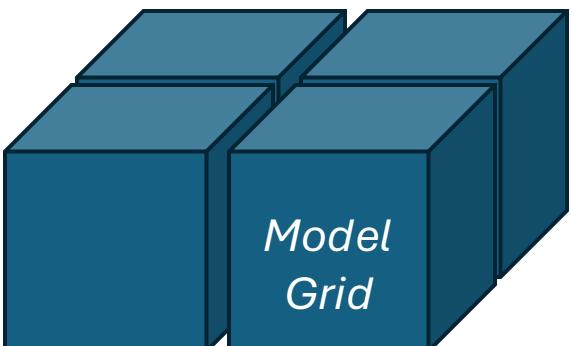
```
In [1]: import xarray as xr # xarray for accessing and manipulating data
```

```
In [2]: path = "/gws/pw/j07/workshop/ARIA_src_data/"  
t_path = path + "VERIFY_eORCA025_MED_UKESM_19900101_20710101_grid_T.nc"
```

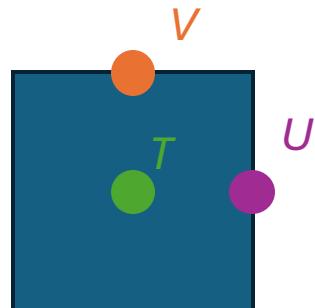
```
In [3]: ds = xr.open_dataset(t_path)
```



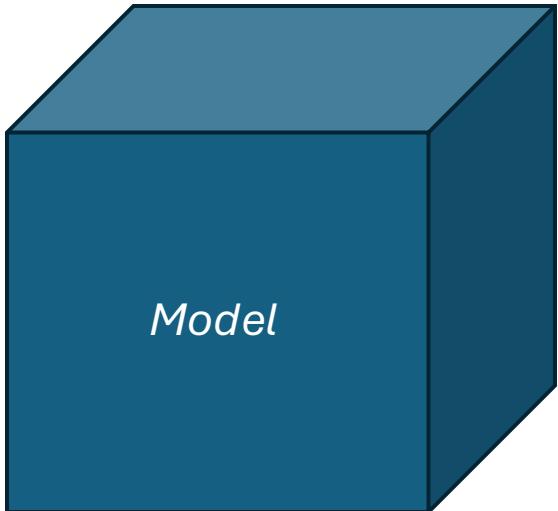
Challenge 1: Introduction



Model
Grid



Model
Cell



T-Grid Data

```
Out[4]: <xarray.Dataset> Size: 7GB
Dimensions:                                (time_counter: 972, y: 400, x: 400, axis_nbounds: 2)
Coordinates:
  nav_lat          (time_counter, y, x) float32 622MB ...
  nav_lon          (time_counter, y, x) float32 622MB ...
  time_centered    (time_counter) object 8kB ...
  * time_counter   (time_counter) object 8kB 1990-01-16 00:00:00 ... 2...
Dimensions without coordinates: y, x, axis_nbounds
Data variables:
  mldr10_1          (time_counter, y, x) float32 622MB ...
  qns_oce           (time_counter, y, x) float32 622MB ...
  qsr_oce           (time_counter, y, x) float32 622MB ...
  sbt              (time_counter, y, x) float32 622MB ...
  sos              (time_counter, y, x) float32 622MB ...
  taum             (time_counter, y, x) float32 622MB ...
  time_centered_bounds (time_counter, axis_nbounds) object 16kB ...
  time_counter_bounds (time_counter, axis_nbounds) object 16kB ...
  tos              (time_counter, y, x) float32 622MB ...
  wfo              (time_counter, y, x) float32 622MB ...
  zos              (time_counter, y, x) float32 622MB ...
Attributes: (12/14)
  name:      MEANS_OUT/eORCA025_MED_UKESM_1m_19900101_19921230_grid_T
  description: ocean T grid variables
  title:     ocean T grid variables
  Conventions: CF-1.6
  timeStamp: 2021-Nov-06 06:11:48 GMT
  uuid:      89941c43-5b52-4f24-8d21-fff8d0897b4c
  ...
  jbegin:    0
  nj:       76
  file_name: eORCA025_MED_UKESM_1m_19900101_19921230_grid_T_199001-1990...
  TimeStamp: 06/11/2021 18:12:14 +0000
  history:   Wed Nov 12 13:39:58 2025: ncks -d x,801,1200 -d y,801,1200 ...
  NCO:       netCDF Operators version 5.3.3 (Homepage = http://nco.sf.ne...
```



Challenge 1: Introduction

- Selecting variables shows the associated meta data
- The example here is for shortwave radiation

```
ds.qsr_oce # we can inspect individual variables

<xarray.DataArray 'qsr_oce' (time_counter: 972, y: 400, x: 400)> Size: 622MB
[155520000 values with dtype=float32]
Coordinates:
    nav_lat      (y, x) float32 640kB ...
    nav_lon      (y, x) float32 640kB ...
    time_centered (time_counter) object 8kB ...
    * time_counter (time_counter) object 8kB 1990-01-16 00:00:00 ... 2070-12-...
Dimensions without coordinates: y, x
Attributes:
    standard_name:      net_downward_shortwave_flux_at_sea_water_surface
    long_name:          solar heat flux at ocean surface
    units:              W/m2
    online_operation:   average
    interval_operation: 1350 s
    interval_write:     1 month
    cell_methods:       time: mean (interval: 1350 s)
```



Challenge 1: Introduction

Horizontal mean, reducing
to 1D time series



```
qsr_mean = ds.qsr_oce.mean(["x","y"])

qsr_mean # the dimensions have been collapsed to time only

<xarray.DataArray 'qsr_oce' (time_counter: 972)> Size: 4kB
array([
  18.476593, 27.610086, 44.851322, 71.09863, 97.54053,
  120.826416, 123.72689, 106.80954, 71.47048, 40.029934,
  22.477606, 15.635142, 17.499926, 25.980236, 44.822742,
  68.59293, 98.48609, 120.40601, 129.7014, 107.867874,
  70.8415, 40.07153, 22.752756, 15.448783, 16.896503,
  24.80583, 42.873085, 70.22497, 96.894714, 121.878204,
  124.35743, 105.791374, 69.4333, 39.713993, 21.998865,
  15.246516, 16.881447, 26.137096, 41.921658, 67.182686,
  99.329956, 125.30074, 127.40543, 107.29668, 70.63323,
  40.50422, 23.103567, 16.292387, 17.892117, 25.44668,
  44.97561, 68.975296, 97.99222, 125.18179, 131.85834,
  111.15676, 71.27184, 40.056988, 22.876837, 15.727722,
  17.68961, 27.04288, 43.43357, 71.9569, 100.051926,
  130.91718, 131.25401, 110.51976, 72.60136, 40.982906,
  22.747236, 15.689405, 17.004135, 27.034433, 44.861935,
  68.21526, 95.91902, 119.6671, 127.46469, 108.8054,
  70.29683, 40.953907, 23.567652, 16.141476, 17.427288,
  27.914274, 44.219173, 67.68102, 96.98384, 126.17667,
  131.86609, 109.63617, 72.350426, 40.107918, 22.094315,
  16.225048, 17.533821, 27.340672, 45.542423, 69.96093,
  ...
])
```

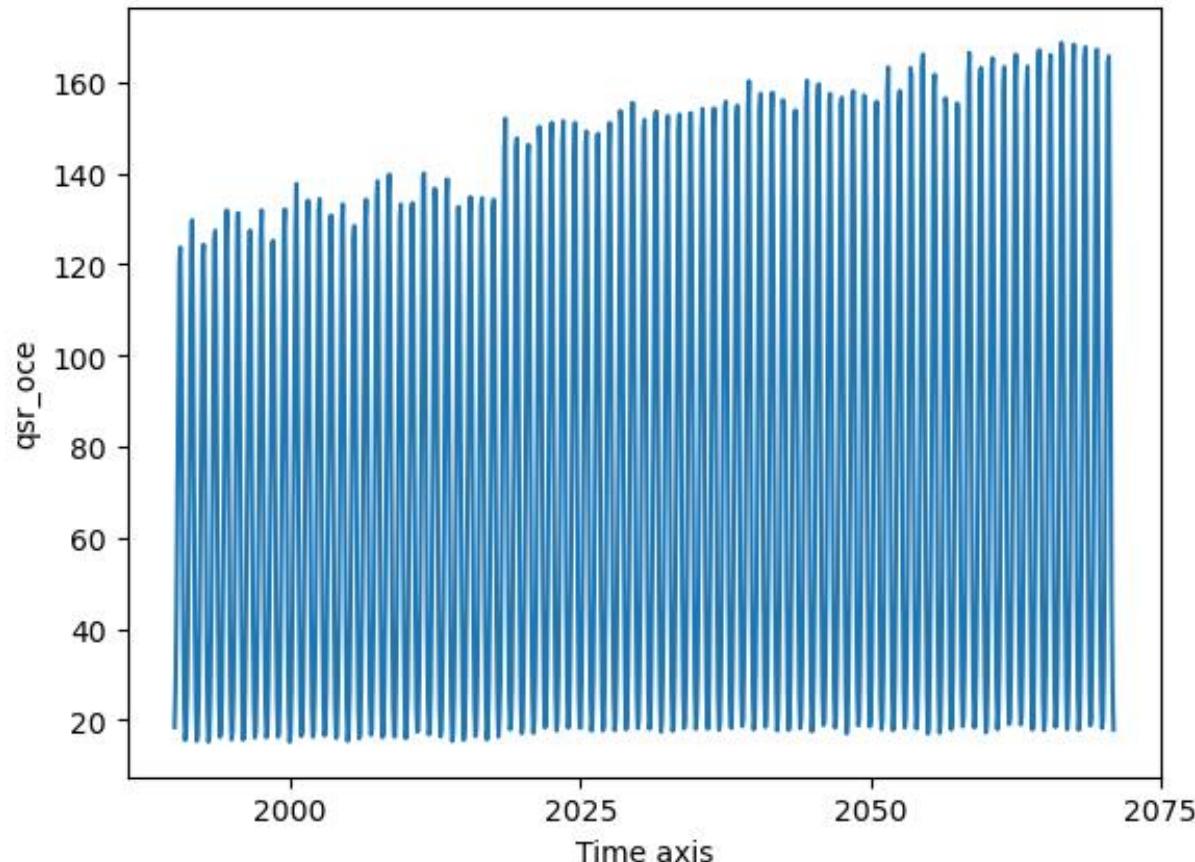


Challenge 1: Introduction

Plotting of timeseries...

```
qsr_mean.plot() # the xarray plot function allows us to quickly plot data
```

```
[<matplotlib.lines.Line2D at 0x7efbdcc5d1df0>]
```

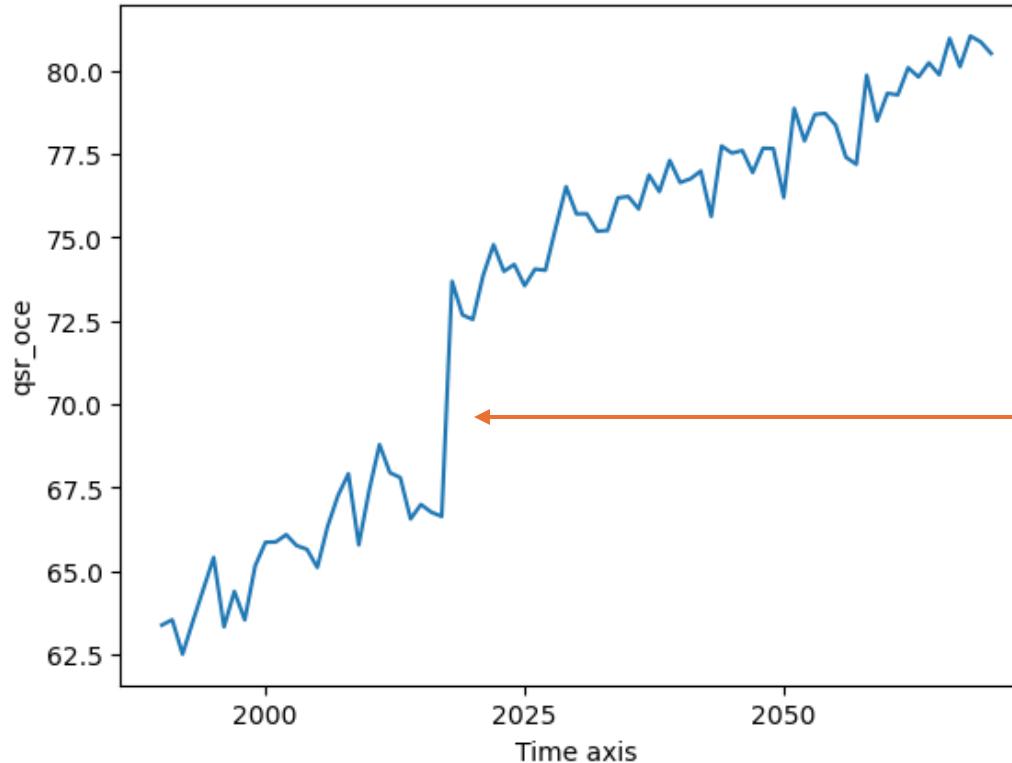


Challenge 1: Introduction

Using resample to extract a yearly mean of the timeseries

```
monthly_qsr = qsr_mean.resample(time_counter="YS").mean()  
monthly_qsr.plot() # interesting
```

[<matplotlib.lines.Line2D at 0x7efbdb981df0>]

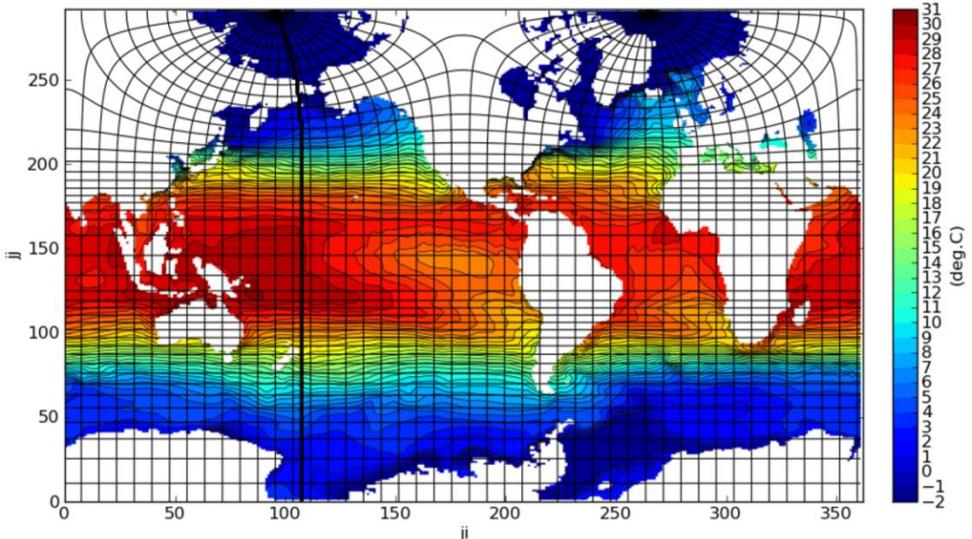


Interesting jump



Challenge 1: Introduction

NEMO simulations commonly use a tri-polar grid, which distorts latitude and longitude.

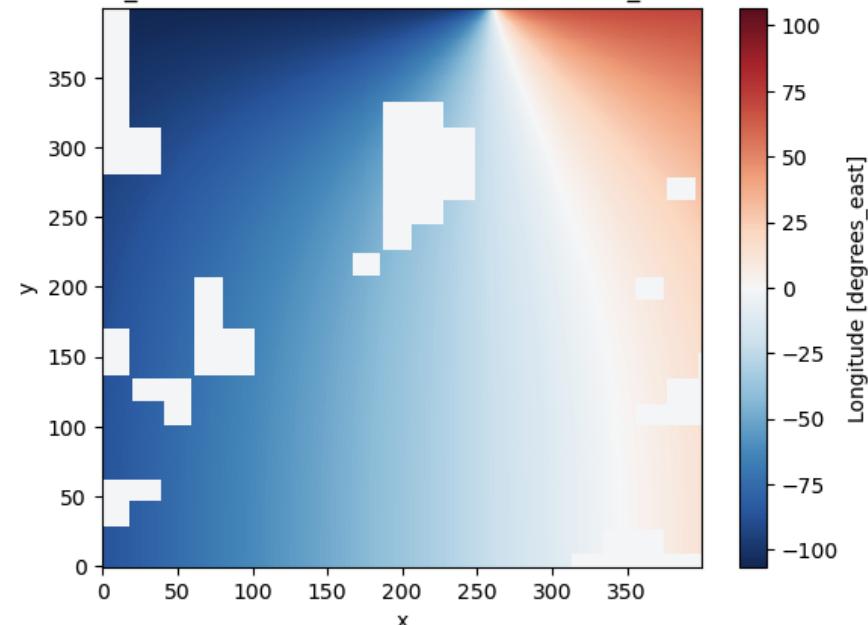


<https://brodeau.github.io/sosie/>

```
ds.nav_lon.plot()
```

```
[10]: <matplotlib.collections.QuadMesh at 0x7f02e95a1dc0>
```

```
time_centered = 1990-02-16 00:00:00, time_count...
```

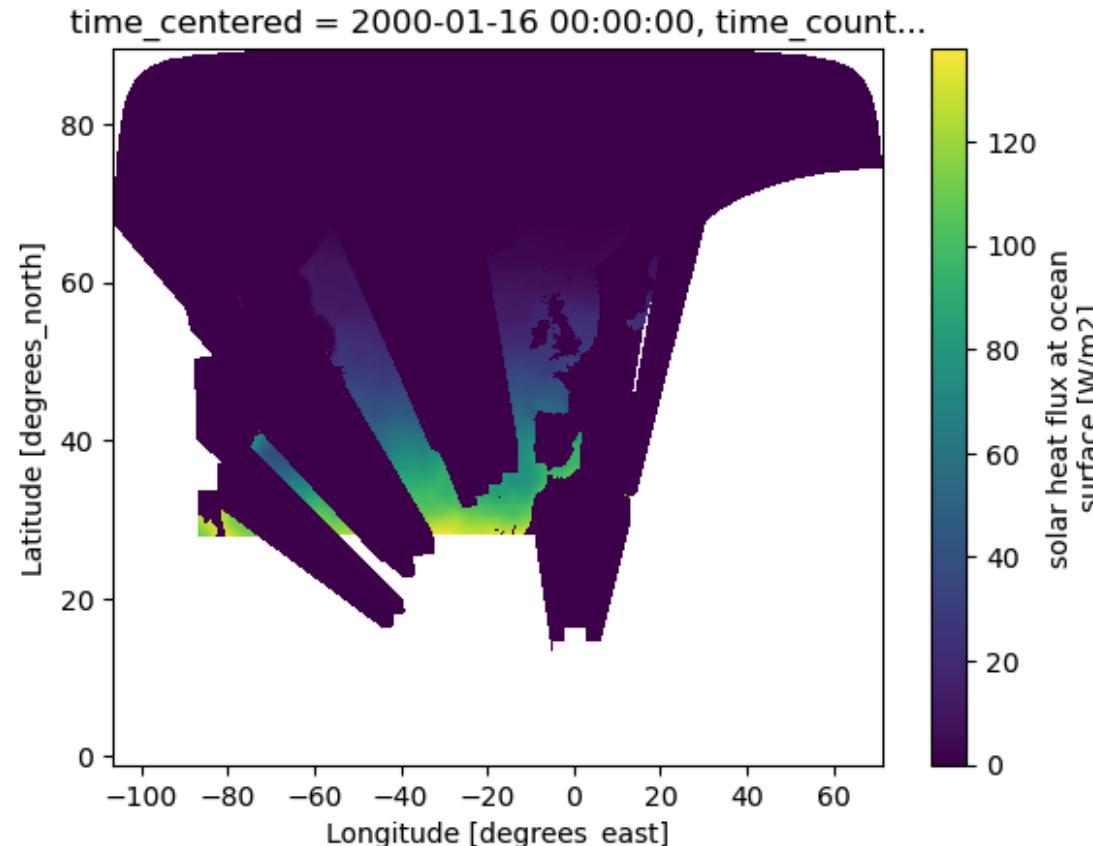


Challenge 1: Introduction

Extracting spatial 2D data

```
[25]: # first we select a date  
surf_rad_y2000m01 = ds.qsr_oce.sel(time_counter="2000-01")
```

```
[27]: # now we can plot...  
surf_rad_y2000m01.plot(x="nav_lon",y="nav_lat")
```



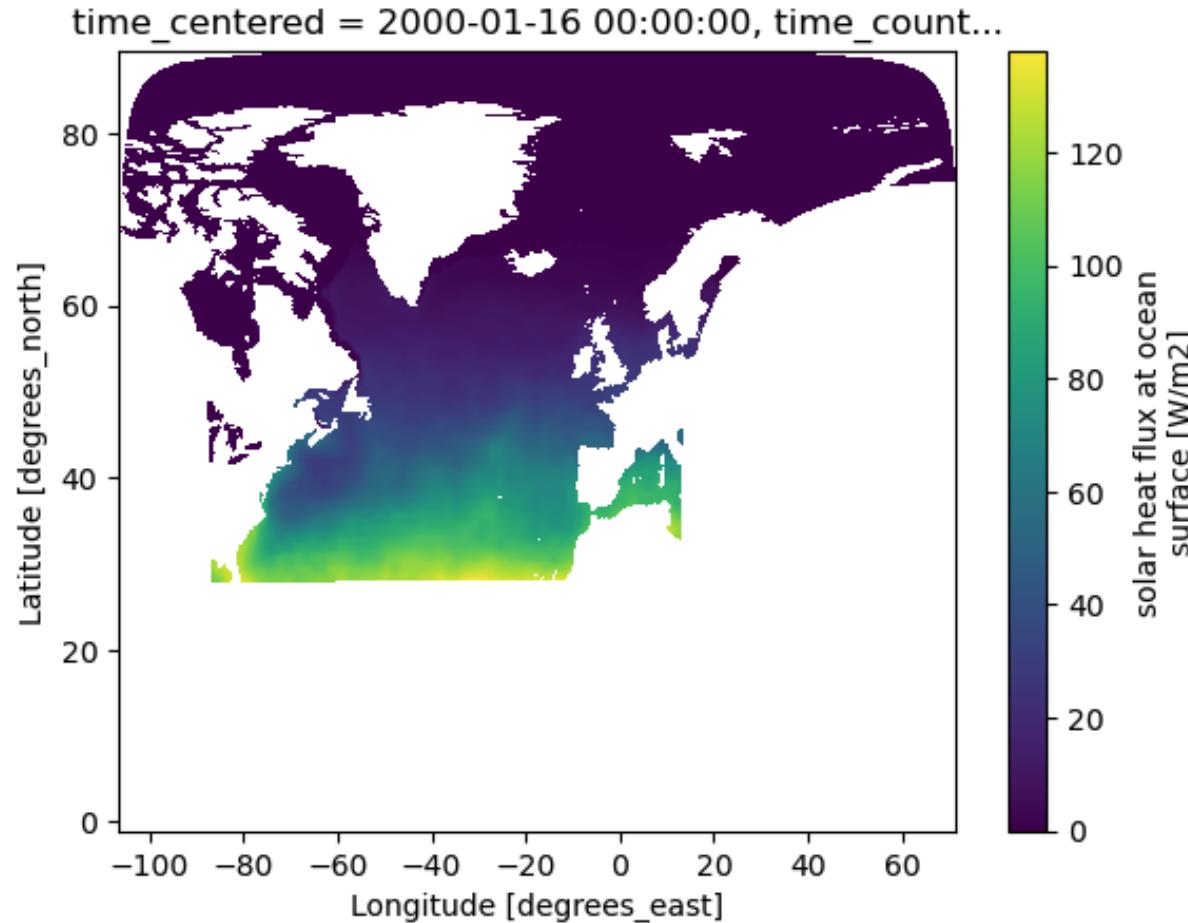
This looks rather peculiar.
Masking of land is needed
to avoid this.



Challenge 1: Introduction

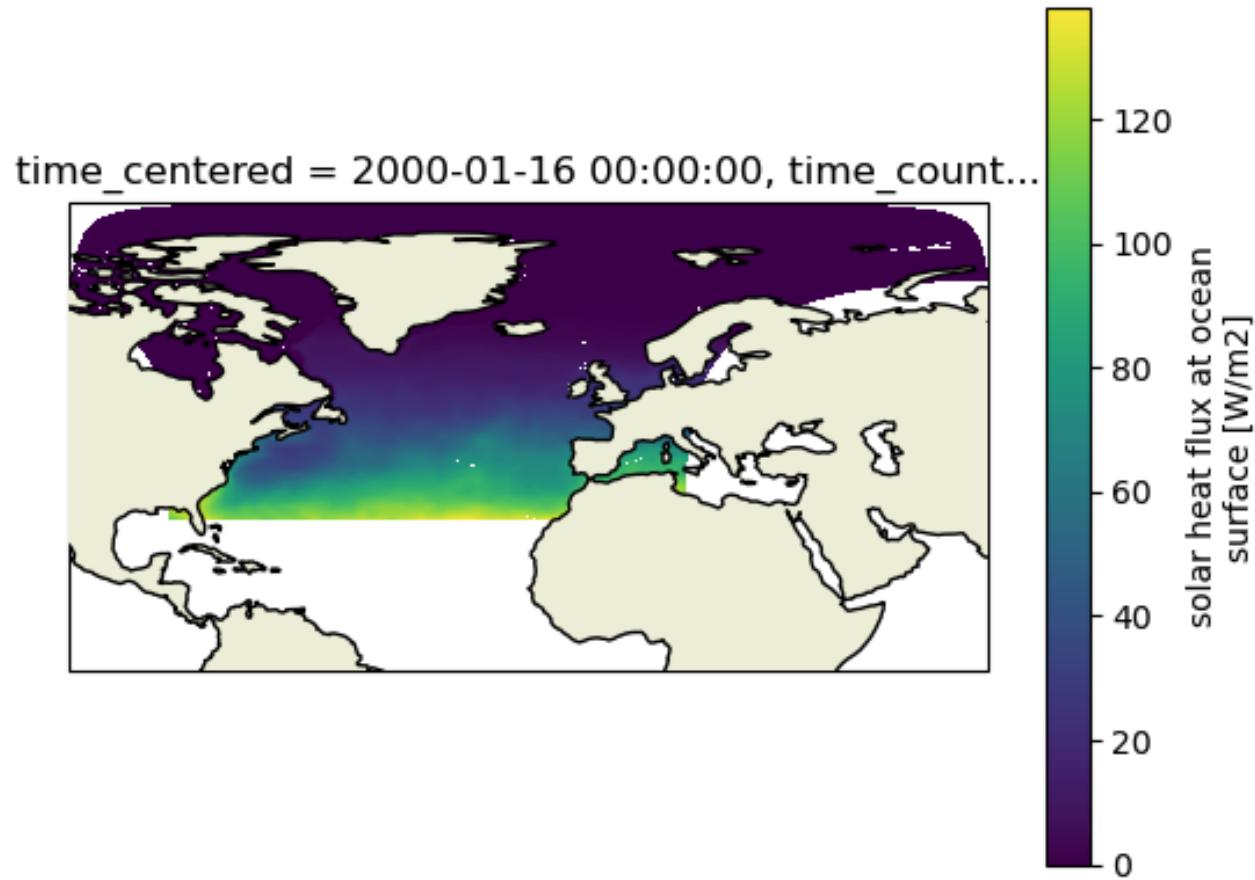
```
[16]: surf_rad_y2000m01 = surf_rad_y2000m01.where(dom_cfg.top_level == 1)
      surf_rad_y2000m01.plot(x="nav_lon",y="nav_lat")
```

Result after masking of land



Challenge 1: Introduction

```
[17]: # now let us plot using a transform  
ax=plt.subplot(projection=ccrs.PlateCarree())  
ax.add_feature(cfeature.LAND, zorder=100, edgecolor='k')  
surf_rad_y2000m01.plot(ax=ax, x="nav_lon",y="nav_lat", transform=ccrs.PlateCarree())
```



We can do fancier projections with cartopy



Challenge 1: Introduction

Groups!

Group 1	Group 2	Group 3	Group 4
Anna Hardisty	Ella Beaven	Jack Wharton	Shivendra Singh Verma
Nina Brendling	Guy Ludford	Ellie Fisher	Evan Hambly
Thanasis Giannakopoulos	Molly Hammond	Ella Tanner	Brad Neimann
Isaac Foreman	Alejandro Coca-Castro	Andrea Quintanilla	Gaby Johnson
Amethyst Eicher	Aoife Ní Bhuachalla	Charli Frisby	Valeria Mascolo



Challenge 1: Introduction

Now over to you...

You have access to:

verify.domcfg.nc

VERIFY_eORCA025_MED_UKESM_19900101_20710101_grid_T.nc

VERIFY_eORCA025_MED_UKESM_19900101_20710101_icemod.nc

VERIFY_eORCA025_MED_UKESM_3D_19900101_20710101_grid_T.nc

VERIFY_eORCA025_MED_UKESM_3D_19900101_20710101_grid_U.nc

VERIFY_eORCA025_MED_UKESM_3D_19900101_20710101_grid_V.nc



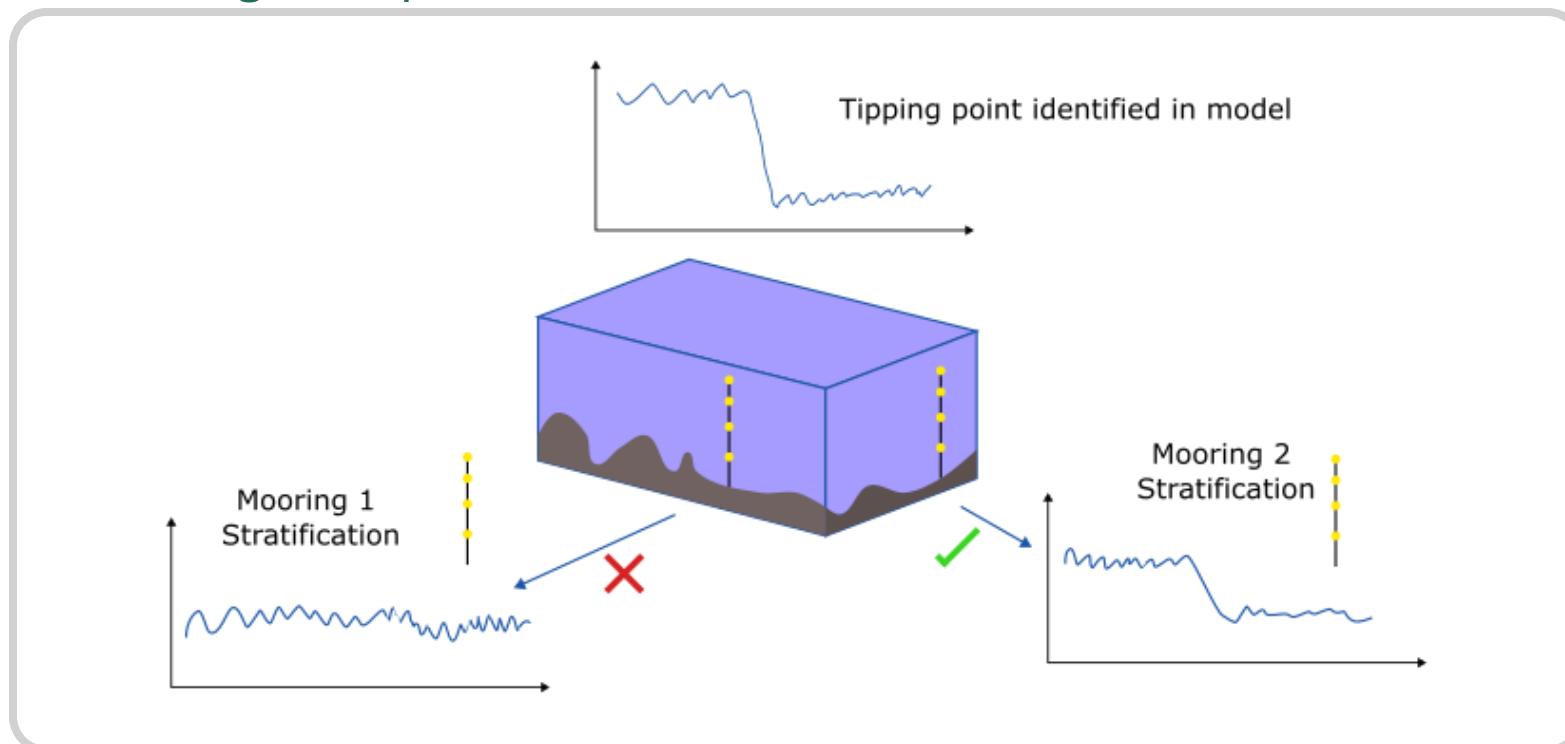
Challenge 2



Challenge 2: Virtual Observations

“Demonstrate how one might devise an observational campaign to capture this tipping point”

Mooring Example



Challenge 2: Virtual Observations

3D dataset

Note the
differences in
variable format



```
# Your dataset path and filename are /gws/pw/j07/workshop/ARIA_src_data/  
# VERIFY_eORCA025_MED_UKESM_3D_19900101_20710101_grid_T.nc  
  
import xarray as xr  
  
path = "/gws/pw/j07/workshop/ARIA_src_data/"  
t_path3d = path + "VERIFY_eORCA025_MED_UKESM_3D_19900101_20710101_grid_T.nc"  
  
ds = xr.open_dataset(t_path3d, chunks={"x":400,"y":400})  
  
ds
```

xarray.Dataset

► Dimensions: (time_counter: 852, deptht: 75, axis_nbounds: 2, y: 400, x: 400)

▼ Coordinates:

deptht	(deptht)	float32 0.5058 1.556 ... 5.902e+03	
nav_lat	(y, x)	float32 dask.array<chunksize=(400, 400...)	
nav_lon	(y, x)	float32 dask.array<chunksize=(400, 400...)	
time_centered	(time_counter)	object dask.array<chunksize=(852,), m...	
time_counter	(time_counter)	object 1990-01-16 00:00:00 ... 2070-12...	

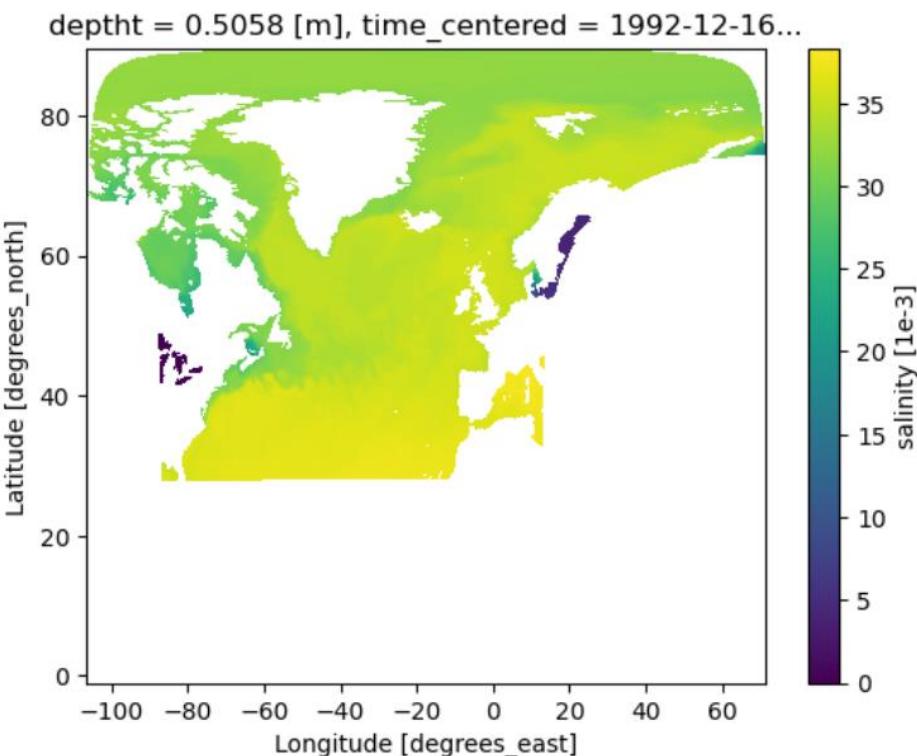


Challenge 2: Virtual Observations

You can change the depth you are looking at.

Take a 2d slice at one time
Eg. surface

Look at it over the 2d field.
This is similar to challenge1.



```
# Select a date and a depth level
# Use the sel() function to select the data based on values of a specific dimension.
# Use the isel() function to select the data based on the index position.

# This is an example for december 1992:
# 3D salinity field for that month: (deptht, y, x)
so_3d = ds.so.sel(time_counter="1992-12")

# 2D field for the same time at a specific depth level (deptht dimension). Change the deptht
Salinity = ds.so.sel(time_counter="1992-12").isel(deptht=0)
# alternatives:
# Salinity = ds.so.sel(time_counter="1992-12").sel(deptht=0, method="nearest") # in this case
# Salinity = ds.so.isel(time_counter=35).isel(deptht=0) # in this case the time is not the dimension

# mask the data to cover the land regions using the domain file.
dom_cfg = xr.open_dataset(path + "verify.domcfg.nc") # open the domain
Salinity = Salinity.where(dom_cfg.top_level == 1) # apply masks

# Finally plot your salinity.
Salinity.plot(x="nav_lon", y="nav_lat")
```



Challenge 2: Virtual Observations

Generate a mask using the salinity field.
(this is a hack – usually the mask would
be generated with the model outputs)



```
# Mask for the full dataset. 4D (time_counter, deptht, y, x)
mask = xr.where(ds.so == 0, float("nan"), 1.0)
```

Select one point in space (lat/lon) and time (time_counter) to look at all depth.
First find the closest model point to your coordinates, then get all depth at one time.

```
# One more recommended way to pick a point is to first select a latitude and longitude you want to investigate,
target_lat = 70
target_lon = 10

# Then find the closest model point to these coordinates
dist = ((ds.nav_lat - target_lat)**2 + (ds.nav_lon - target_lon)**2)**0.5

# Find the index of the closest point by:
dist_1d = dist.stack(points=("y", "x")).load() # Turning the 2D grid (y, x) into a 1D list of points called "point"
ip = dist_1d.argmin("points") # getting index of the point with minimum distance (closest point)

# Recover the original (y, x) indices from the stacked MultiIndex
j = dist_1d["y"].isel(points=ip).item()
i = dist_1d["x"].isel(points=ip).item()

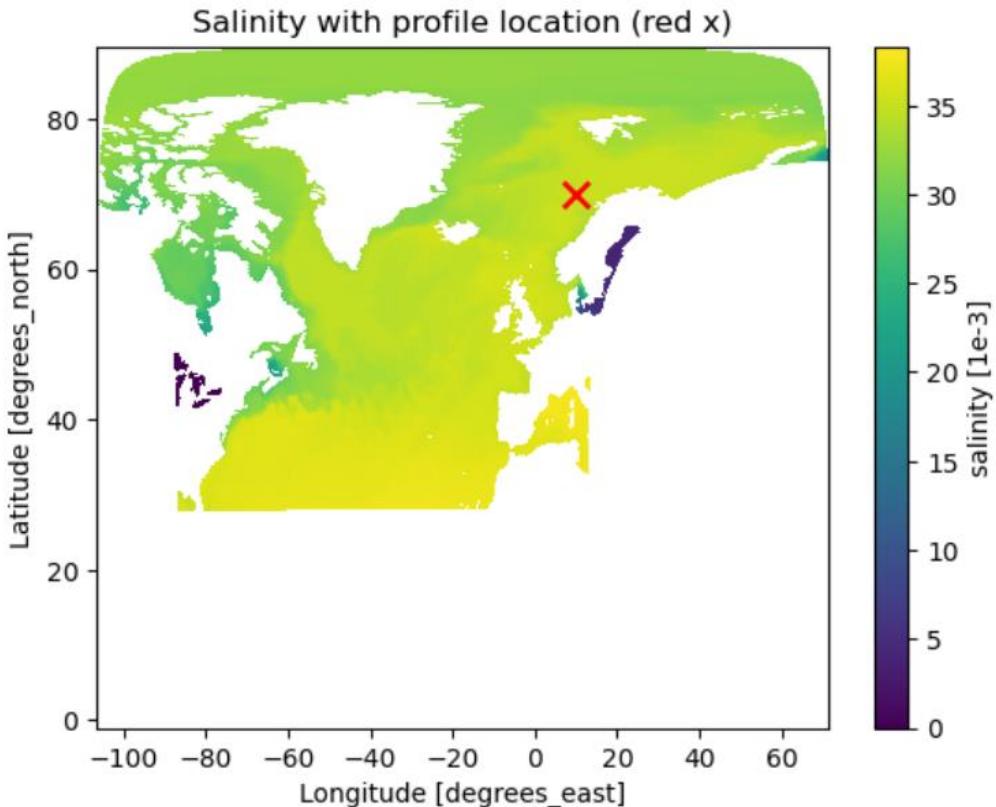
# Use (j, i) with isel to get the vertical profile (all depths, one time)
salinity_profile = ds.so.sel(time_counter="1992-12").isel(y=j, x=i)

#print(salinity_profile)
#print(salinity_profile.deptht.values)
```



Challenge 2: Virtual Observations

Check the point is where you wanted it



```
import matplotlib.pyplot as plt

# Lon/Lat of the chosen point
pt_lon = ds.nav_lon.isel(y=j, x=i)
pt_lat = ds.nav_lat.isel(y=j, x=i)

# Plot salinity and overlay the red 'x'
fig, ax = plt.subplots()

Salinity.plot(x="nav_lon", y="nav_lat", ax=ax)
ax.plot(pt_lon, pt_lat, "rx", markersize=10, mew=2)

ax.set_title("Salinity with profile location (red x)")
```



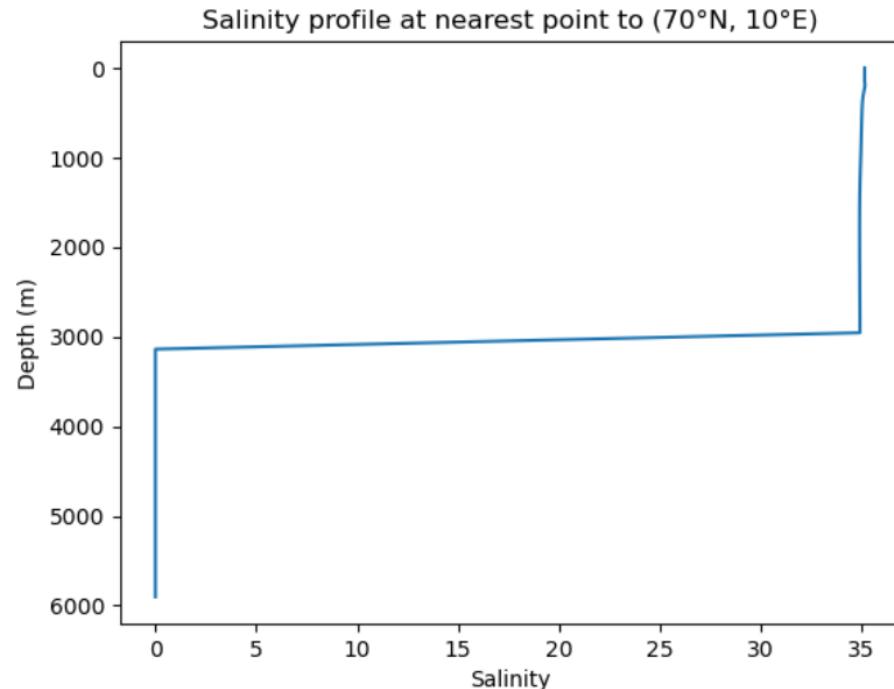
Challenge 2: Virtual Observations

Plot your salinity profile un-masked



```
# Simple profile: salinity vs depth
# import matplotlib.pyplot as plt

salinity_profile.plot(y="deptht")      # or x="deptht" if you prefer depth on x-axis
plt.gca().invert_yaxis()               # optional: depth increasing downward
plt.title("Salinity profile at nearest point to (70°N, 10°E)")
plt.xlabel("Salinity")
plt.ylabel("Depth (m)")
```



Challenge 2: Virtual Observations

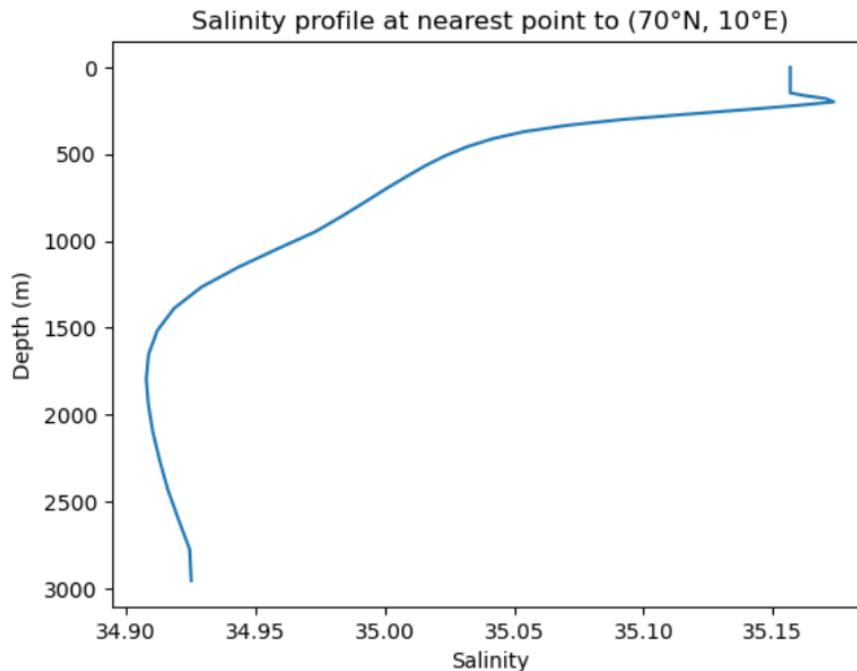
Plot your salinity profile masked



```
# Extract the 1D mask profile at the same time and point
mask_profile = mask.sel(time_counter="1992-12").isel(y=j, x=i)

# Apply the mask: keep values where mask == 1, set others to NaN
salinity_profile_masked = salinity_profile.where(mask_profile == 1)

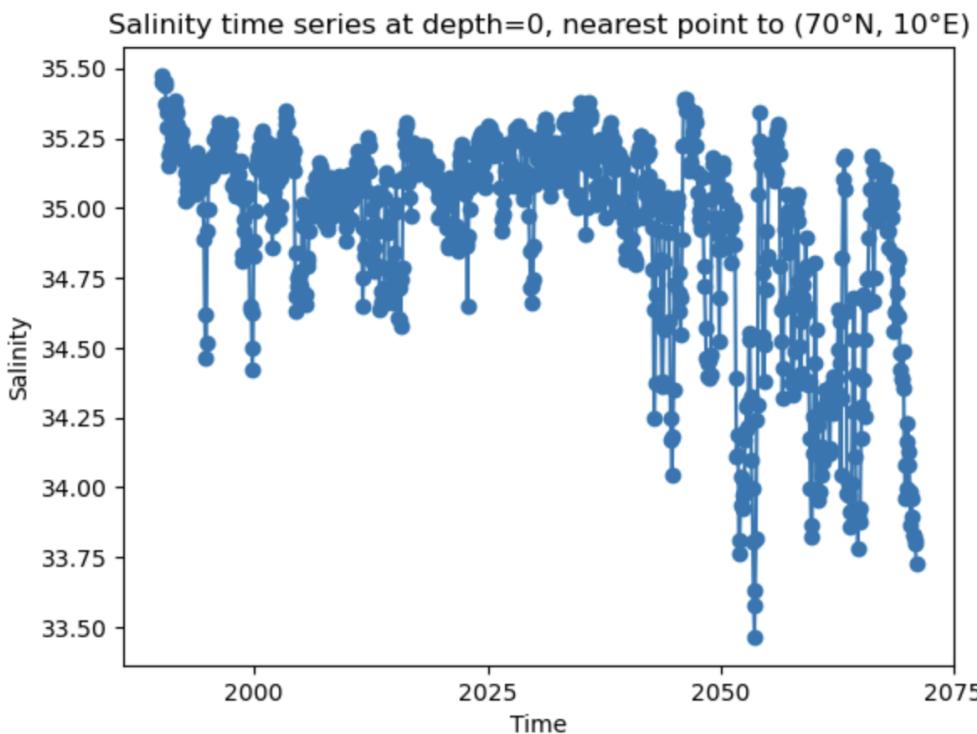
# Plot masked profile
salinity_profile_masked.plot(y="deptht")    # or x="deptht"
plt.gca().invert_yaxis()
plt.title("Salinity profile at nearest point to (70°N, 10°E)")
plt.xlabel("Salinity")
plt.ylabel("Depth (m)")
```



Challenge 2: Virtual Observations

Try a timeseries:

Same point,
one depth,
all times.



```
# Time series at surface (depth=0) for the selected grid point (j, i)
salinity_ts = ds.so.isel(depth=0, y=j, x=i) # dims: time_counter

# plot the timeseries
import matplotlib.pyplot as plt
fig, ax = plt.subplots()

# Line + scatter time series
salinity_ts.plot(ax=ax)
ax.scatter(salinity_ts["time_counter"].values,
           salinity_ts.values)

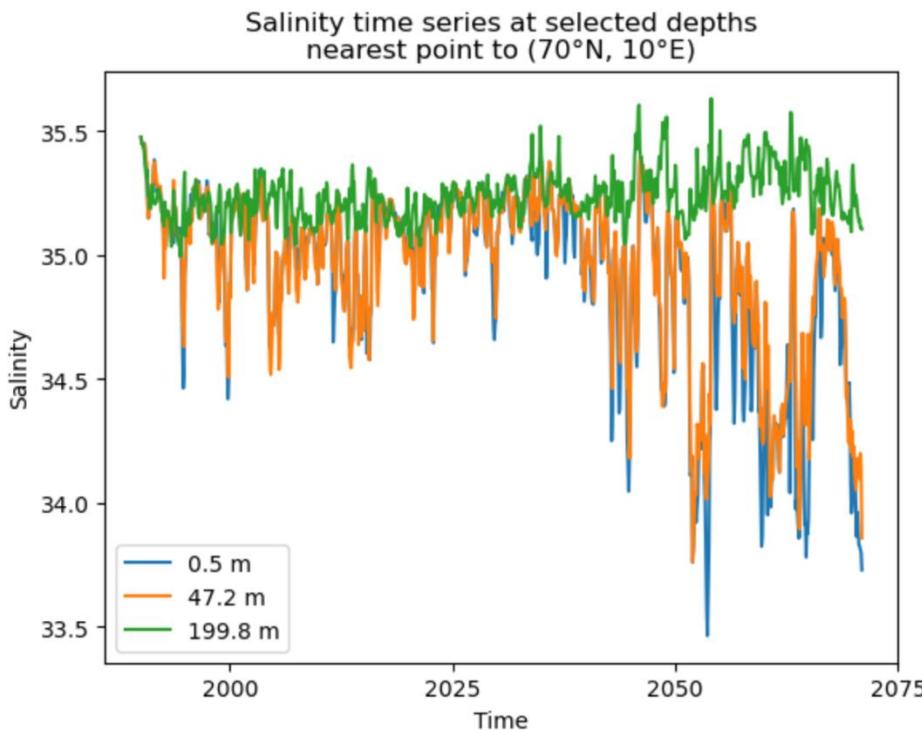
ax.set_title("Salinity time series at depth=0, nearest point to (70°N, 10°E)")
ax.set_xlabel("Time")
ax.set_ylabel("Salinity")

plt.show()
```



Challenge 2: Virtual Observations

You can do
timeseries at
different depths



```
# Hand-picked depths in metres
depths = [0, 50, 200]    # change as you like

fig, ax = plt.subplots()
for d in depths:
    # time series at depth ~ d, for point (j, i)
    salinity_ts = ds.so.sel(deptht=d, method="nearest").isel(y=j, x=i) # find the closest
    label = f"{float(salinity_ts.deptht.values):.1f} m"      # Label with the actual depth
    salinity_ts.plot(ax=ax, label=label)

ax.set_title("Salinity time series at selected depths\nnearest point to (70°N, 10°E)")
ax.set_xlabel("Time")
ax.set_ylabel("Salinity")
ax.legend()
plt.show()
```



Challenge 2: Virtual Observations

But a better way to look at one profile over multiple times is a *Hovmöller plot*.

First select all depth at one point for the full timeseries subset and mask it.



```
: # Extract salinity and mask at the chosen horizontal point (j, i)
salinity_hov = ds.so.isel(y=j, x=i)    # dims: ('time_counter', 'deptht')
mask_hov     = mask.isel(y=j, x=i)      # same dims. This uses the 4D mask you have generated at

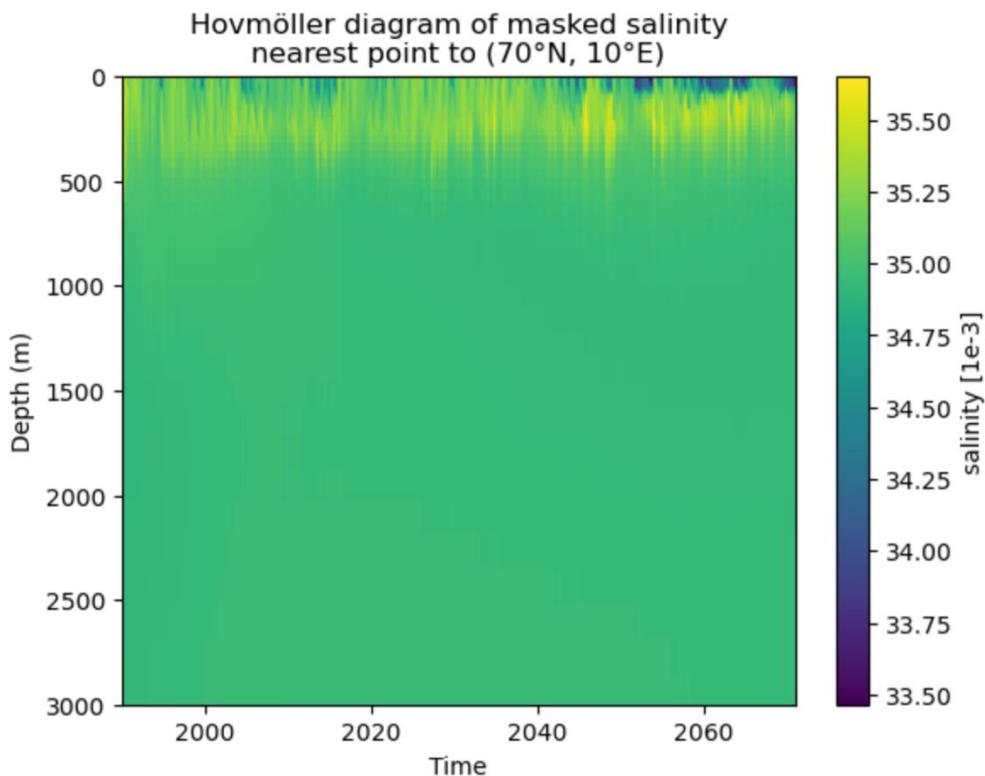
# Apply mask to your salinity slice.
salinity_hov_masked = salinity_hov.where(mask_hov == 1)

# (optional sanity check)
print(salinity_hov_masked.dims)        # should be ('time_counter', 'deptht')
```



Challenge 2: Virtual Observations

Then plot your *Hovmöller plot*. →



```
# Hovmöller plot: time vs depth
fig, ax = plt.subplots()

salinity_hov_masked.plot(
    x="time_counter",
    y="deptht",
    ax=ax,
)
# x-limits: first and last time (you can change these)
t0 = salinity_hov_masked["time_counter"].values[0]
t1 = salinity_hov_masked["time_counter"].values[-1]
ax.set_xlim(t0, t1)
# y-limits: e.g. top 500 m (and flip so 0 is at the top)
ax.set_ylim(3000, 0)

ax.set_title("Hovmöller diagram of masked salinity\nnearest point to (70°N, 10°E)")
ax.set_xlabel("Time")
ax.set_ylabel("Depth (m)")
```



Challenge 2: Virtual Observations

Now over to you...

Play around, try different points, and different subsets of data...

Can you see a tipping point?

What about testing different types of observations? Ship-based underway or saildrones sampling perhaps?

Autosub



Ship-based
underway



Unmanned surface
vehicles



Challenge 3: Advanced



Challenge 3: Tipping Point detection

What we will cover

- We will introduce methods for tipping point detection
- Key metrics include variance and auto-correlation

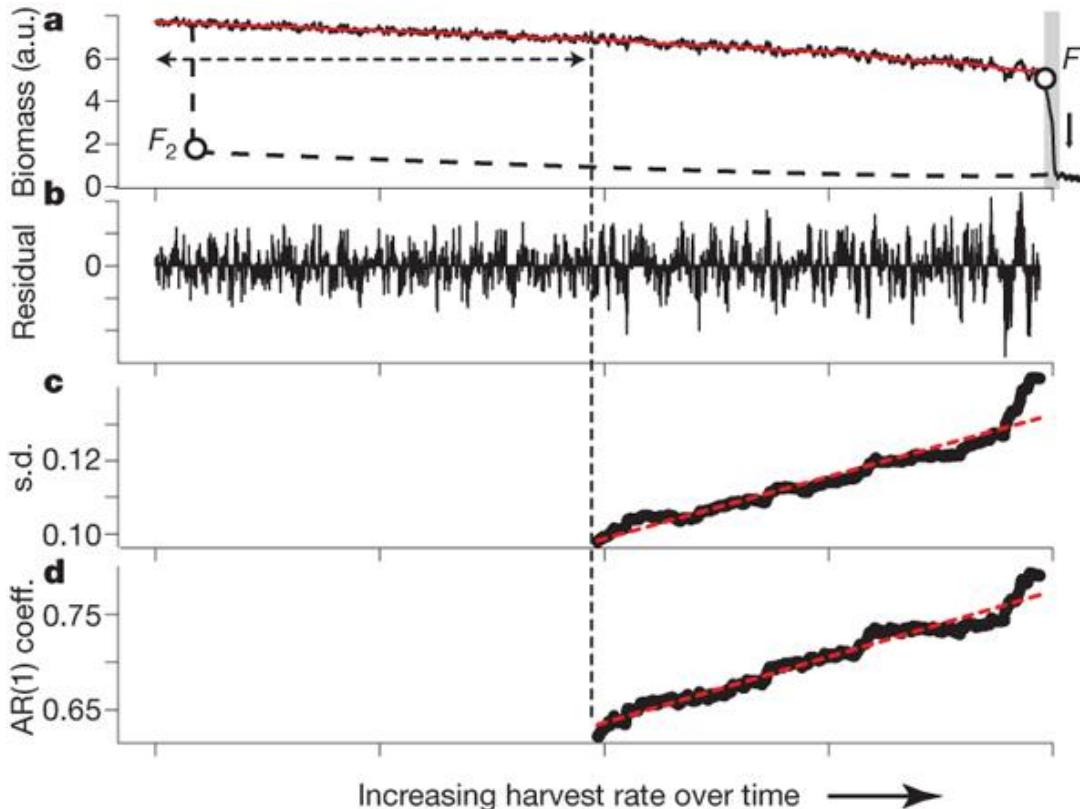


Challenge 3: Tipping Point detection

Tipping Point Overview

- Systems can behave differently in the lead up to a tipping point
- Statistical analysis can reveal the changes

1. Raw timeseries



2. Detrended timeseries

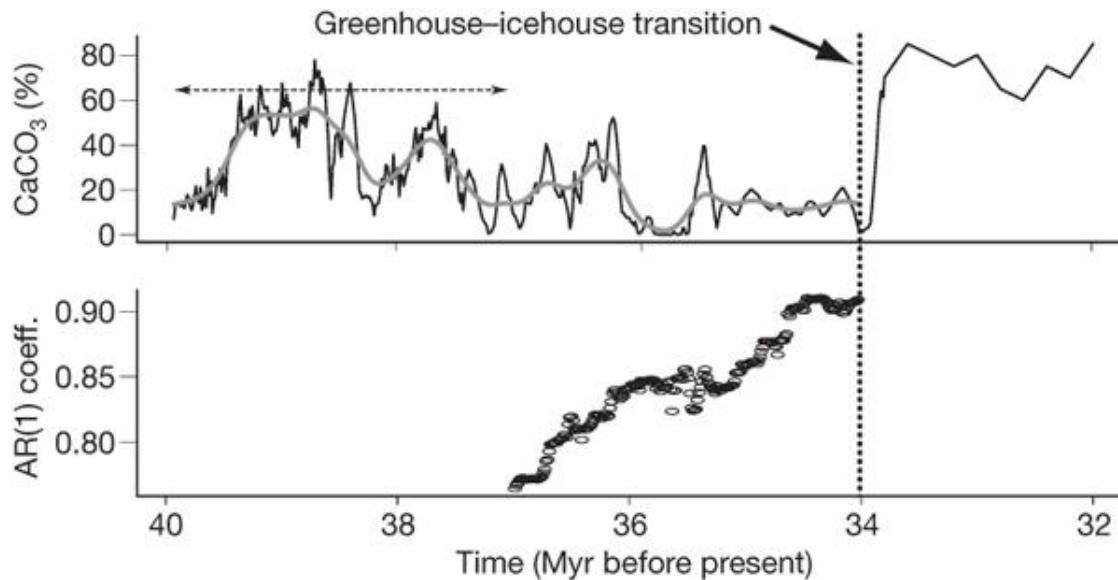
3. Variance^{1/2}

4. Lag-1 Autocorrelation

Challenge 3: Tipping Point detection

Example in palaeo record

- We observe an abrupt transition in the percentage of CaCO_3
- Associated rise on auto-correlation



Challenge 3: Tipping Point detection

Worked example

- Example with surface temperature extraction
- Extract a location as before

```
path = "/gws/pw/j07/workshop/ARIA_src_data/"
t_path = path + "VERIFY_eORCA025_MED_UKESM_19900101_20710101_grid_T.nc"

tos = xr.open_dataset(t_path).tos # get surface temperature data

# Extract a time series but choosing a latitude and longitude you want to investigate
target_lat = 70
target_lon = 10

# Then find the closest model point to these coordinates
dist = ((tos.nav_lat - target_lat)**2 + (tos.nav_lon - target_lon)**2)**0.5

# Find the index of the closest point by:
dist_1d = dist.stack(points=("y", "x")) # Turning the 2D grid (y, x) into a 1D list
ip = dist_1d.argmin("points") # getting index of the point with minimum distance (closest)

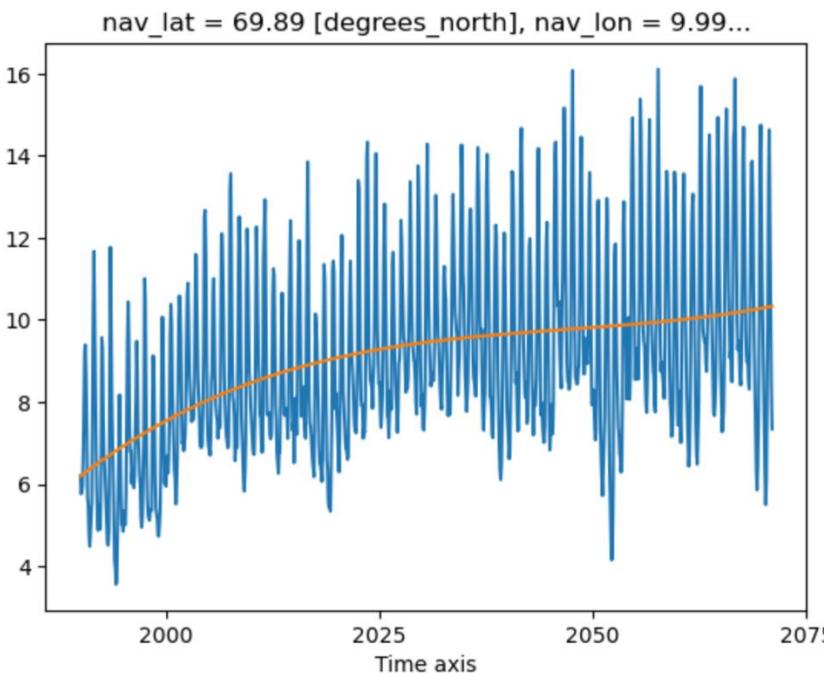
# Recover the original (y, x) indices from the stacked MultiIndex
j = dist_1d["y"].isel(points=ip).item()
i = dist_1d["x"].isel(points=ip).item()

tos_timeseries = tos.isel(x=i,y=j)
```



Challenge 3: Tipping Point detection

We then detrend the data by fitting a curve to the timeseries and removing the fitted “trend”



```
# This function is designed to remove the trend from the timeseries
def detrend_dim(da, dim, deg=1):
    # detrend along a single dimension
    p = da.polyfit(dim=dim, deg=deg)
    fit = xr.polyval(da[dim], p.polyfit_coefficients)
    return da - fit, fit

# detrended_tos is the detrended temperature timeseries and fit defined the fitted curve
detrended_tos, fit = detrend_dim(tos_timeseries, "time_counter", 3)

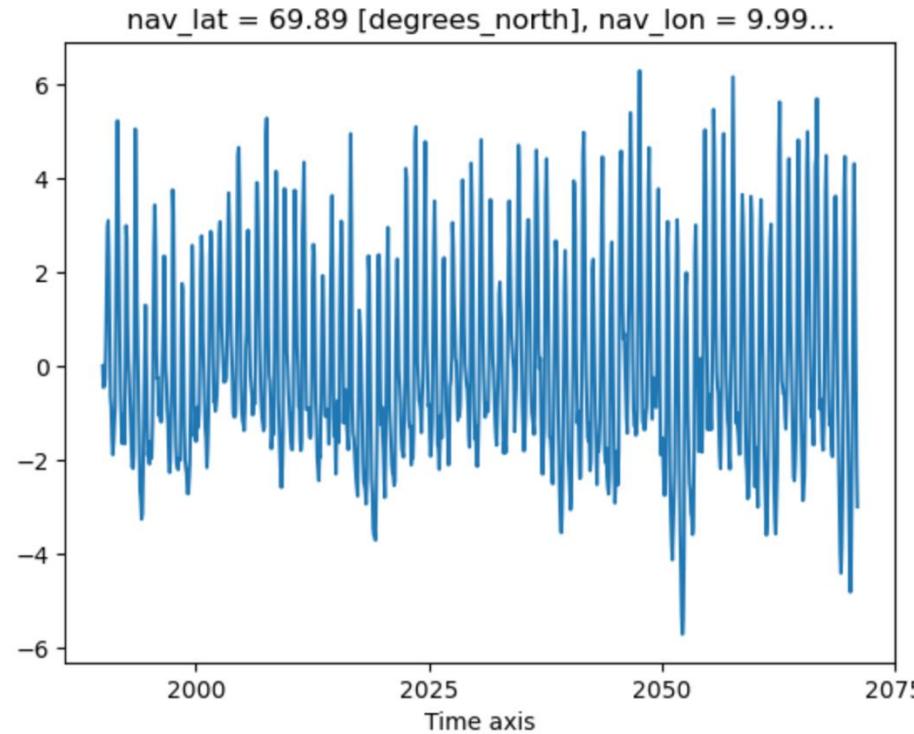
# we can show the fit here
tos_timeseries.plot()
fit.plot()
```



Challenge 3: Tipping Point detection

Detrended result

```
: # once detrended, we recover temporal variability  
detrended_tos.plot()
```

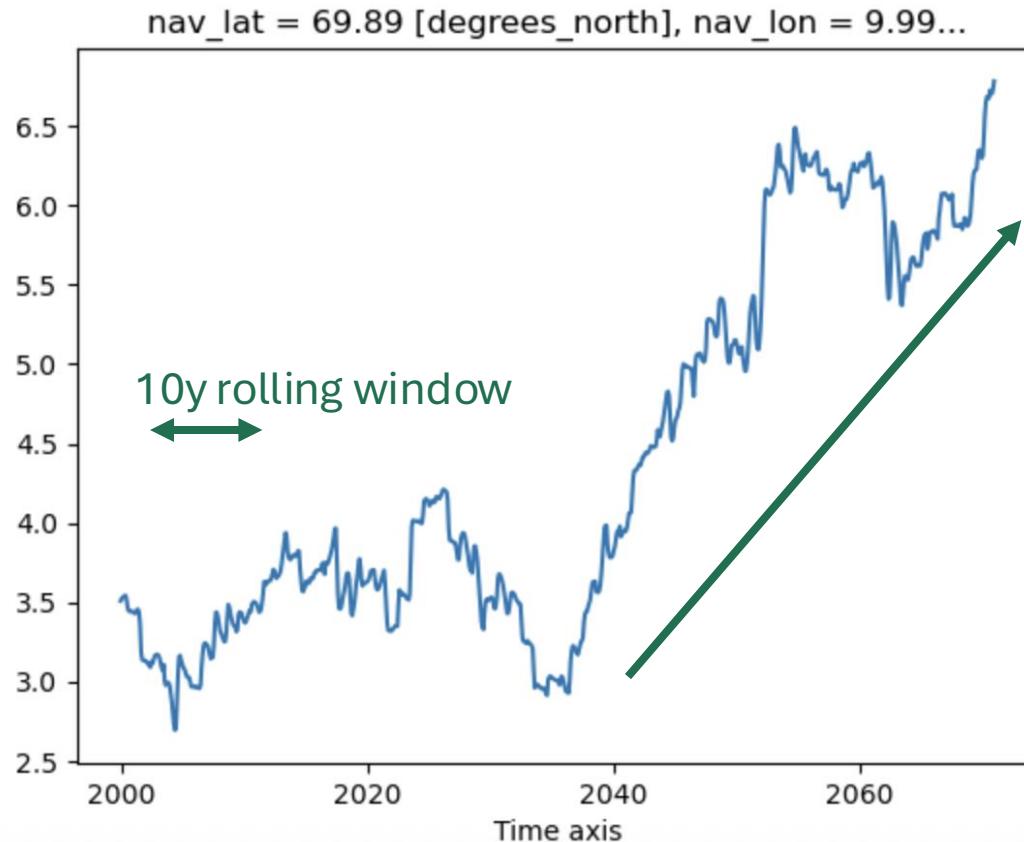


Challenge 3: Tipping Point detection

We then find the variance over a rolling window – ten years in this case

```
# let's calculate variance over a rolling window
tos_variance = detrended_tos.rolling(time_counter=120).var()

# We then plot the variance to evaluate the emergence of a tipping point in this system
tos_variance.plot()
```



Are we heading for a tipping point!?



Challenge 3: Tipping Point detection

Over to you

- Use the notebooks as a guide (see additional autocorrelation example)
- Try out different timeseries data. What about Challenge 1 data?
- Can you think of other methods that might be useful for detecting tipping?

